



Develop an Extendable REST API server and Composite UI Based on Microservices

Pitt Capstone Project Fall 2019

NetApp Software-Defined Core Infrastructure
September 13, 2019



Agenda

1) Introductions

- Name
- Interests
- Previous experience

2) A little about NetApp

- What do we do?
- What is Software Defined Core Infrastructure?

3) The Project

4) Introduction to concepts

- REST
- Containers & Docker

5) Scheduling

- Working together with us

About NetApp

- Data Authority in the Hybrid Cloud
 - Software, systems, and services for data
 - Deploying and building providers IT environment
- \$6 billion+ in revenue
- Fortune 500® Company
- More than 10,000 employees in 150 offices worldwide
- Our office: PTC – Pittsburgh Technical Center
 - Acquired Spinnaker Networks in 2003
 - 200 Employees – almost all engineering



Introductions

NetApp Engineers
Pitt Students

Our Backgrounds

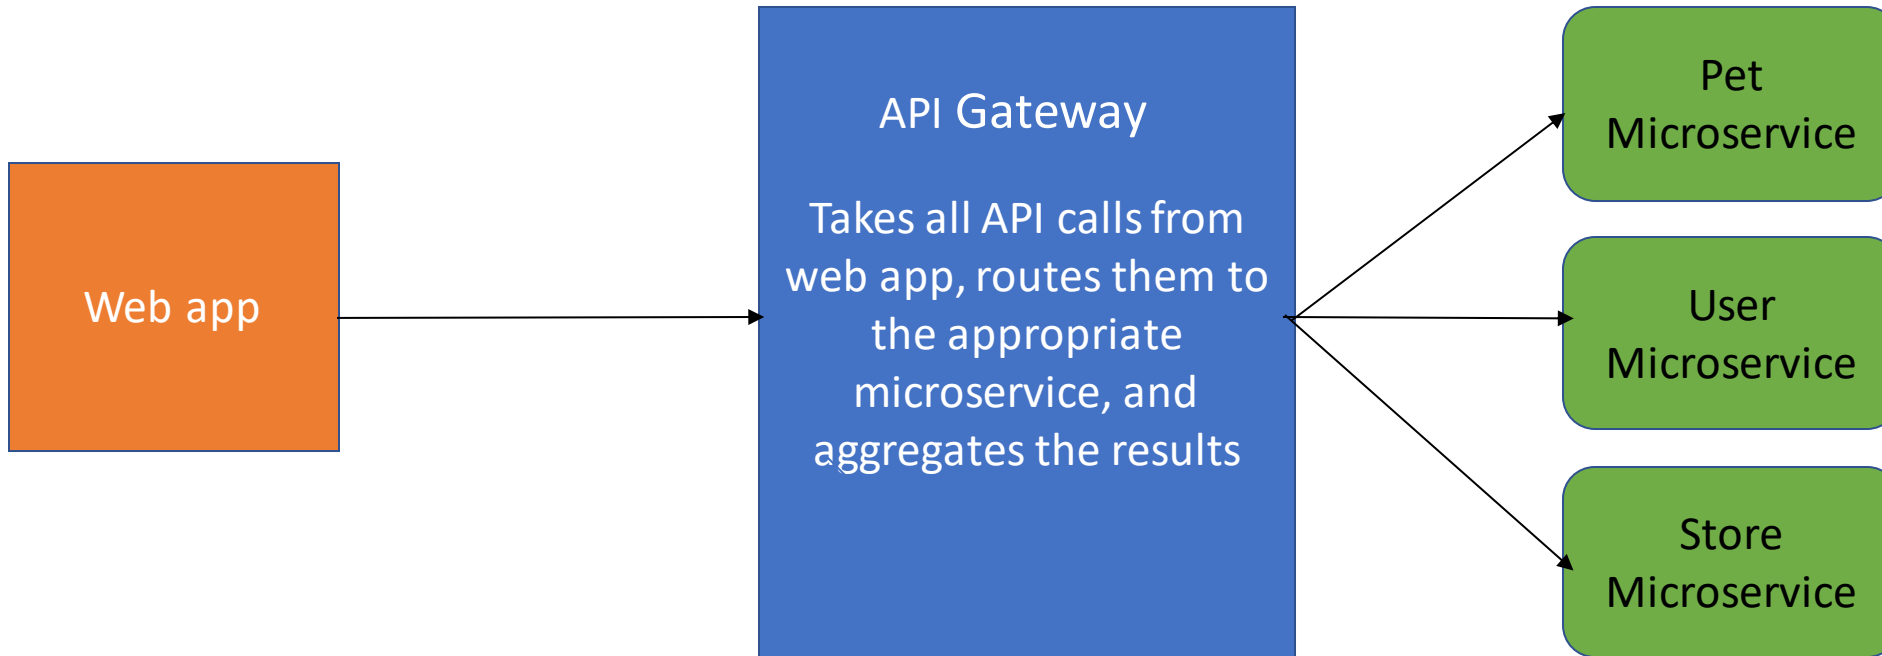
Where we were before NetApp



- **Charlie Mietzner** (Charlie.Mietzner@netapp.com)
 - Graduated in Fall 2016 from Pitt with BS in CS and BA in English Literature
 - Capstone Project for NetApp in Fall 2016
 - Started at NetApp in mid-July 2017
- **Twesha Mitra** (Twesha.Mitra@netapp.com)
 - Graduated in Spring 2018 with BS in Economics and BS in CS.
 - Full time at NetApp since June 2018

The Project

A proof-of-concept Web application built on microservices and an API gateway



Why a petstore?

It's not about the destination; it's about the journey.

- A Web Application for the theoretical owner and employees of a pet store.
- Keeps track of pets, customers, and customer orders in the store.
- Under the covers it must be implemented with a microservice architecture
- An investigative and experimental project
- There isn't a "right answer"
- Future groups can model their work after your architecture
- We are equally interested in the pitfalls encountered

Deliverables

A proof-of-concept Web application built on microservices and an API gateway

Deliverables/Requirements:

- Workflows should be working as expected
- All endpoints/actions not covered in workflows should also be supported
- Web GUI (client app) that implements the workflows

Workflow 1:

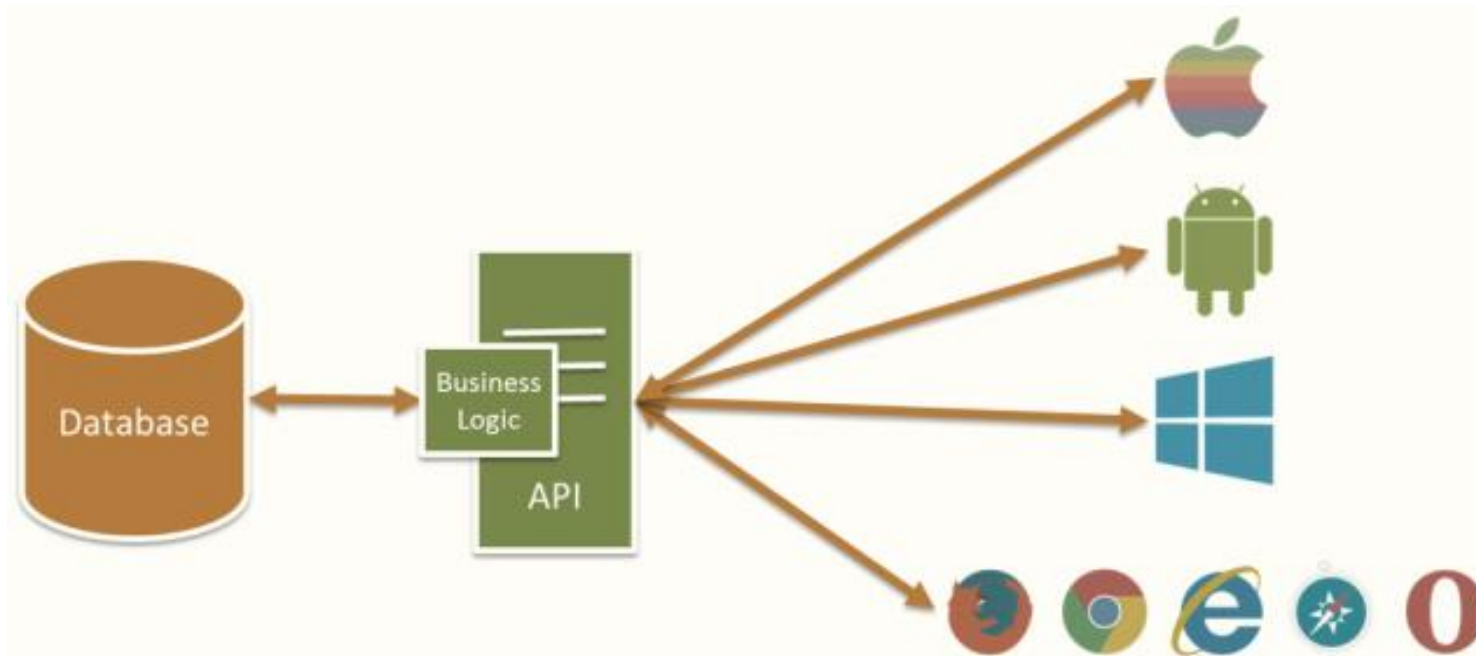
- Register a customer
- View all available pets
- Customer buys a pet
- Owner creates a store order for the pet
- Owner updates pet/user info

Workflow 2:

- Existing customer wants to update their contact information
- Customer has a pet of type A, want to buy another pet of type A
- Create store order for the pet and update relevant information

What is REST (REpresentational State Transfer)

- An application should interact with a resource with only the following knowledge:
 - Resource identifier (URLs, also known as endpoint)
 - Action (GET, POST, PUT, PATCH, DELETE)
 - Understanding of the representation returned (JSON, XML, HTML)



REST constraints

- Stateless
 - Each request from client to server must contain all information necessary to understand the request
 - Scalable and distributable
- Cacheable
 - Resources can be cached, but they must declare themselves cacheable
- Layered system
 - Each component cannot see beyond the immediate layer with which they are interacting
- Uniform interface
 - Uniform way of interacting with a server independent of device/application
 - Identification & manipulation
- Client-server separation
 - Client application and server application should be able to evolve separately without any dependencies on each other

Containers?

Provides a similar benefit of a Virtual Machine

- A virtual machine uses emulated hardware to run a guest OS; a container is effectively an emulated OS.
- Relative to a VM, containers provide isolation and portability in a more lightweight...well...container

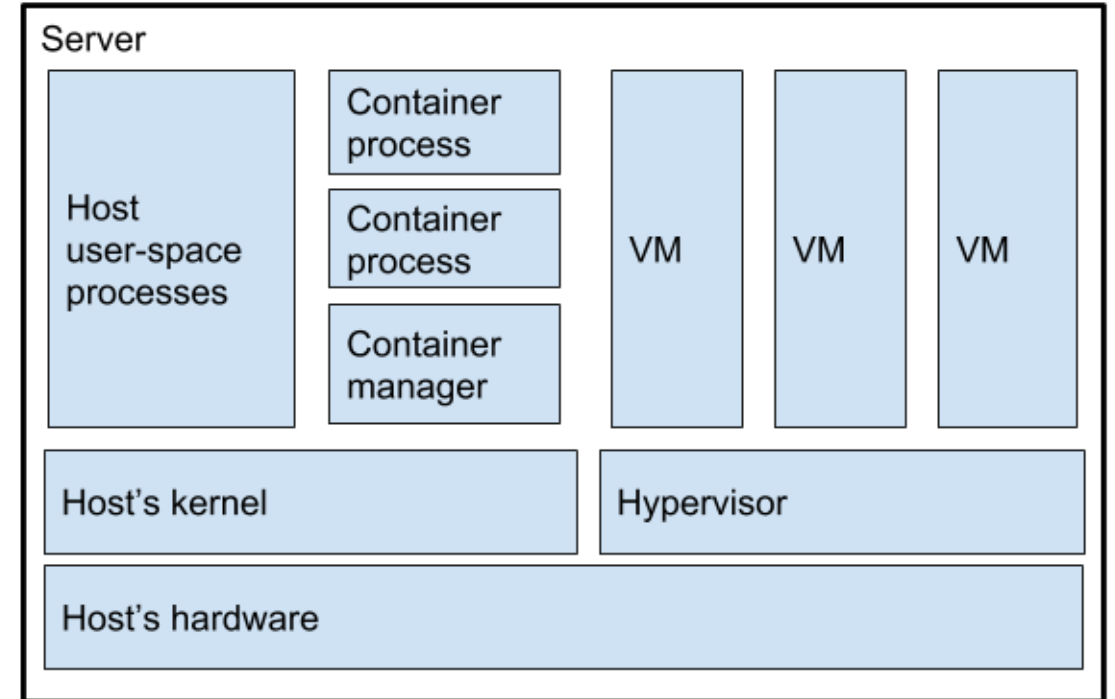


Figure 1. VM vs Container Comparison. Adapted from "What is the difference between a process, a container, and a VM?" By Jessica G. *Medium*. 2019. Retrieved from <https://medium.com/@jessgreb01/what-is-the-difference-between-a-process-a-container-and-a-vm-f36ba0f8a8f7>

Microservices?

A Software Architectural Style

- A software application composed of independently deployable services.
- Think "highly cohesive, loosely coupled".
- Each piece of the application can be developed and deployed independently.
- Each microservice within the application can use different technology stacks, languages, communication protocols—it just needs to satisfy a well-defined API
- Fault isolation. A bug in one microservice doesn't *technically* affect other microservices in the application

API Gateway?

Why do we need one?

- Lots of microservices requires that the client keep track of multiple hosts / service instances.
- Each microservice (by design) could have different interfaces (HTTP, AMQP, RPC, etc...). The client would need to switch between these to interact with the full application.
- Interacting with several microservices requires the client to make several network requests.
- The client would need to implement their own solutions to achieve complicated application-level workflows requiring several microservices.

What should it provide?

- Think big-picture workflows.
- Think about the needs of the consumer of your application.
- There doesn't need to be just *one* API gateway. E.g. Netflix

Scheduling

- Meet once a week with entire team
 - In person every other week
 - Over the phone every other week
- Midterm presentation
- Final Presentation
- Communication
 - Email, Slack
- Progress tracking and source control
 - Github
 - [Creating a project board on github](#)

Where to get started?

- Populate kanban board on github with tasks below
- Schedule midterm/final presentations
- Decide on weekly meeting time & place
- Get a basic REST server up and running (HTTP) in a Docker container
- Come up with a system design
 - Roadmap on how to implement
 - Technologies? Languages?
 - Populate product backlog in github